

Model	KV/LWW	Mem0	Summary	Patch	CAMP
Granite 4 (350M)	46.03 / 13.49	77.84 / 25.78	75.26 / 42.45	81.38 / 43.92	77.68 / 44.92
Qwen3.5 (0.8B)	46.03 / 13.49	85.01 / 25.50	82.49 / 46.51	84.40 / 52.93	81.15 / 50.78
Granite 4 (1B)	46.03 / 13.49	71.90 / 25.35	83.16 / 46.69	82.86 / 50.23	84.22 / 49.22
Llama-3.2 (1B)	46.03 / 13.49	83.06 / 25.12	79.11 / 47.74	81.64 / 46.49	83.96 / 51.75
Qwen3.5 (2B)	46.03 / 13.49	80.77 / 26.33	79.47 / 47.51	83.75 / 52.28	84.01 / 54.00
Llama-3.2 (3B)	46.03 / 13.49	83.38 / 26.23	81.84 / 49.59	85.05 / 54.97	81.78 / 53.31
Mean	46.03 / 13.49	80.33 / 25.72	80.22 / 46.75	83.18 / 50.14	82.13 / 50.66

Table 1: Memory-update and final-memory quality on CAMPBench. Each cell reports Update F1 / Final-memory F1 (%; higher is better), using the same models and checkpoints as the main performance table. Evaluation covers all 24 test scenarios; Update F1 uses the fixed 2,364-turn sample (394 UPDATE and 1,970 NO.OP turns). Final-memory F1 is averaged over scenarios, and Mean over the six models. Bold marks the highest value for each metric in each row. Mem0 denotes the one-pass variant; CAMP uses $k = 5$. No Memory and Recent-window are omitted because these metrics are undefined for them. KV/LWW maintains memory over all original turns and evaluates Update F1 on the same fixed sample.