



Quality-latency trade-off under update load

Base · U40 · U80


Granite 350M




Qwen 0.8B




Llama 1B



Granite 1B



Qwen 2B



Method

- Patch
- k=2
- k=5
- k=10

Update workload

- Base
- U40
- U80

Trajectory: Base → U80

Accuracy budget

- ≥ Patch
- ≤1 pp loss
- ≤3 pp loss