



Quality-latency trade-off under update load

Base · U20 · U40 · U60 · U80


Granite 350M




Qwen 0.8B




Llama 1B




Granite 1B



Qwen 2B



Llama 3B



← Faster · Latency vs. Patch (%)

← Faster · Latency vs. Patch (%)

← Faster · Latency vs. Patch (%)

Method
Patch k=2 k=5 k=10

Update load
Base U20 U40 U60 U80

Accuracy budget
≥ Patch ≤1 pp loss ≤3 pp loss